

# Computer Architecture and Operating System

## Input-Output Systems

Prof. Indranil Sengupta

Department of Computer Science and Engineering

IIT Kharagpur

## Introduction

- Interfacing input/output devices is more complex as compared to interfacing memory systems.
- Why?
  - Wide variety of peripherals (keyboard, mouse, disk, camera, printer, scanner, etc.).
  - Widely varying speeds.
  - Data transfer rate can be regular or irregular.
  - Sizes of data blocks transferred at a time varies widely (few bytes to Kbytes).
- Slower than processor and memory.

## Magnetic Disk (Hard Disk)

- Magnetic disks constitute a traditional method for non-volatile storage of information using magnetic technology.
- Broadly three types of devices appeared:
  - a) **Floppy disk** : made of bendable plastic
  - b) **Magnetic drum** : made of solid metal
  - c) **Hard disk** : made of metal or glass
- All of these rely on a rotating platter (metal or glass or plastic) coated with a thin magnetic material, and use a moveable read/write head to read and write data from / to the disk.
  - Data stored as tiny magnets.

26/03/24

Computer Architecture and Operating System (CS31702)

3



Magnetic drum  
(62.5 Kbytes)



8" floppy disk  
(360 Kbytes)



3.5" floppy disk  
(1.2 Mbytes)



3.5" magnetic disk  
(1 Tbytes)



1.8" solid-state disk  
(512 Gbytes)


26/03/24

Computer Architecture and Operating System (CS31702)

4

## Organization of Data on a Hard Disk



- The hard disks consists of a collection of *platters* (typically, 1 to 5), which are connected together and can spin in unison.
  - Each platter has two recording surfaces, and comes in various sizes (1 – 8 inches).
  - The stack of platter typically rotates at a speed of 5,400 to 10,000 rpm.
  - Each disk surface is divided into concentric circles called *tracks*.
    - The number of tracks per surface can vary from 1000 to 5000.
  - Each track is divided into a number of *sectors* (64 – 200 sectors/track).
    - Typical sector size: 512 – 2048 bytes.
    - Sector is the smallest unit that can be read or written.
- The disk heads for all the surfaces are connected and move together.
- All the tracks under the heads at a given time on all surfaces is called a *cylinder*.



26/03/24

Computer Architecture and Operating System (CS31702)

5



26/03/24

Computer Architecture and Operating System (CS31702)

6

## Disk Access Time

- There are three components to the access time in hard disk:

**a) Seek time:**

- The time required to move the head to the desired track.
- Average seek times are in the range 8 – 20 msec.
- Actual average can be 25 – 30% less than this number, since accesses to disks are often localized.

**b) Rotational delay:**

- Once the head is on the correct track, we must wait for the desired sector to rotate under the head.
- The average delay or latency is the time for half the rotation.
- Examples:
  - For 3600 rpm, average rotational delay =  $0.5 \text{ rotation} / 3600 \text{ rpm} = 8.30 \text{ msec}$
  - For 5400 rpm, average rotational delay =  $0.5 \text{ rotation} / 5400 \text{ rpm} = 5.53 \text{ msec}$
  - For 7200 rpm, average rotational delay =  $0.5 \text{ rotation} / 7200 \text{ rpm} = 4.15 \text{ msec}$

**c) Transfer time:**

- The total time to transfer a block of data (typically, a sector).
- Transfer rates are typically 15 MB/sec or more.
- Transfer time depends on:
  - Sector size
  - Rotation speed of the disk
  - Recording density on the tracks

**Example**

Consider a disk with sector size 512 bytes, 2000 tracks per surface, 64 sectors per track, three double-sided platters, and average seek time of 10 msec.

- a) What is the capacity of the disk?
- b) If the disk platters rotate at 7200 rpm, and one track of data can be transferred per revolution, what is the maximum data transfer rate?

## Data Transfer Techniques


### Data Transfer Techniques

1) **Programmed**: CPU executes a program that transfers data between I/O device and memory.

- a) Synchronous
- b) Asynchronous
- c) Interrupt-driven

*Generally slower as the execution of a program is involved.*

2) **Direct Memory Access (DMA)**: An external controller directly transfers data between I/O device and memory without CPU intervention.



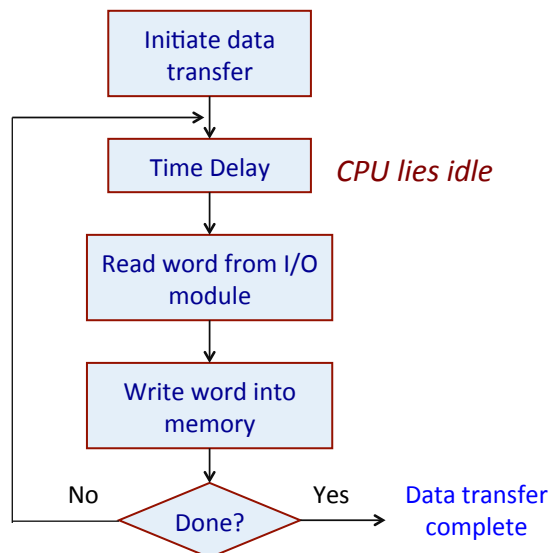
## (a) Synchronous Data Transfer

- The I/O device transfers data at a *fixed rate* that is known to the CPU.
- The CPU initiates the I/O operation and transfers successive bytes/words after giving fixed time delays.
- **Characteristics:**
  - During the time delay, CPU lies idle.
  - Not many I/O devices have strictly synchronous data transfer characteristics.
- A flowchart for synchronous data transfer from an input device is shown on the next slide.

26/03/24

Computer Architecture and Operating System (CS31702)

13



- Error may occur if the input device and the processor get out of synchronization.
- Large number of words cannot be transferred in one go.
- Speed of data transfer depends not only on the speed of I/O device and memory, but also on the execution time of the code.

26/03/24

Computer Architecture and Operating System (CS31702)

14

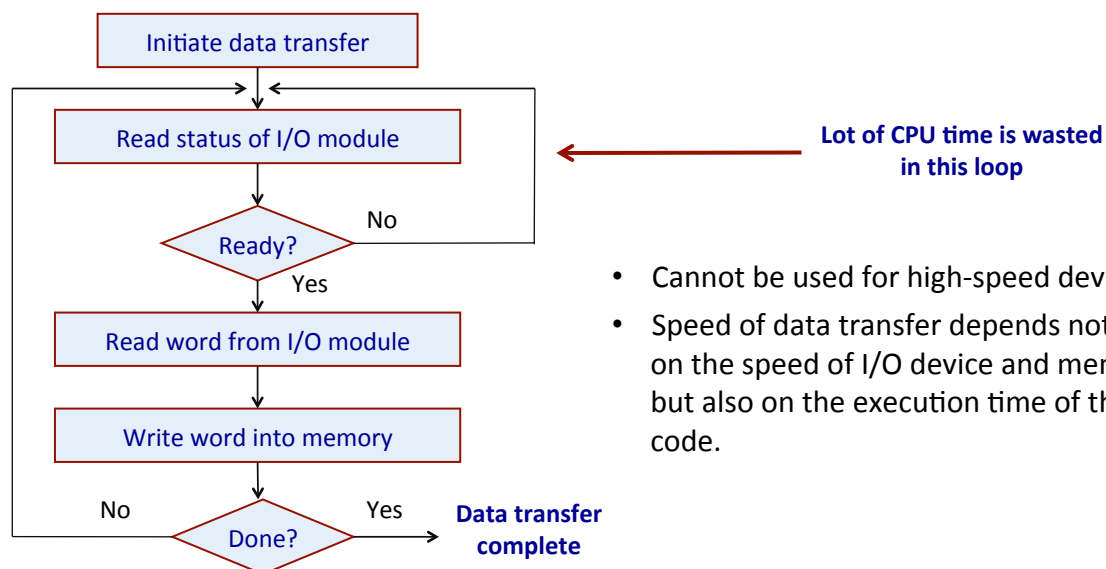
## (b) Asynchronous Data Transfer

- The CPU does not know when the I/O module will be ready to transfer the next word.
- CPU has to check the status of the I/O module to know when the device is ready to transfer the next word.
  - Called *handshaking*.
- Characteristics:
  - While the CPU is checking whether the I/O module is ready, it cannot do anything else.
  - Wasteful of CPU time for slow devices like keyboard or mouse.

26/03/24

Computer Architecture and Operating System (CS31702)

15



- Cannot be used for high-speed devices.
- Speed of data transfer depends not only on the speed of I/O device and memory, but also on the execution time of the code.

26/03/24

Computer Architecture and Operating System (CS31702)

16



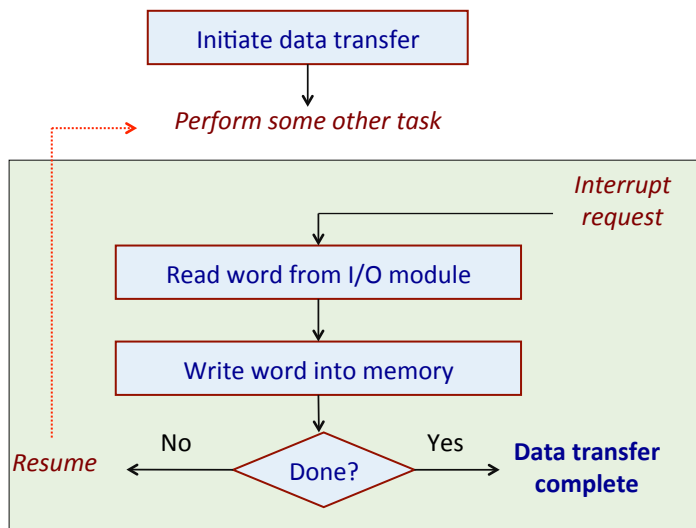
## (c) Interrupt-Driven Data Transfer

- The CPU initiates the data transfer and proceeds to perform some other task.
- When the I/O module is ready for data transfer, it informs the CPU by activating a signal (called *interrupt request*).
- The CPU suspends the task it was doing, services the request (that is, carries out the data transfer), and returns back to the task it was doing.
- Characteristics:
  - CPU time is not wasted while checking the status of the I/O module.
  - CPU time is required only during data transfer, plus some overheads for transferring and returning control.

26/03/24

Computer Architecture and Operating System (CS31702)

17



*This part of the program that gets activated when interrupt request comes is called **interrupt handler (IH)** or **interrupt service subroutine (ISS)**.*


26/03/24

Computer Architecture and Operating System (CS31702)

18

- Typical interfaces for interrupt-driven data transfer:

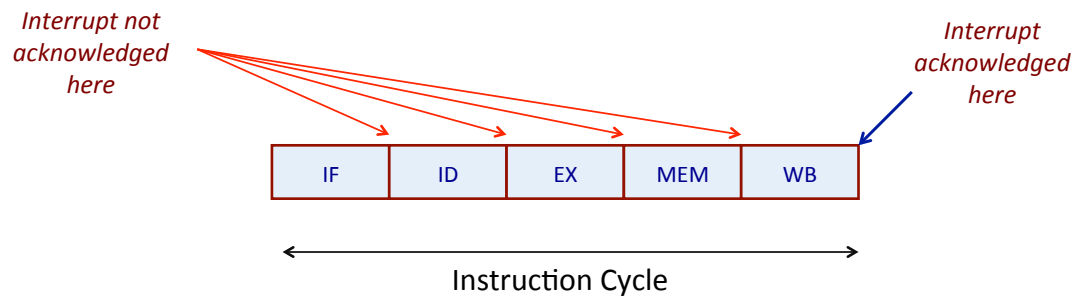
- a) **Interrupt request (INTR)** – External device raises this signal when it wants to draw attention of the CPU.
- b) **Interrupt acknowledge (INTA)** – CPU sends this signal back to the external device to indicate that it has received the INTR signal.



## What happens when an interrupt request arrives?

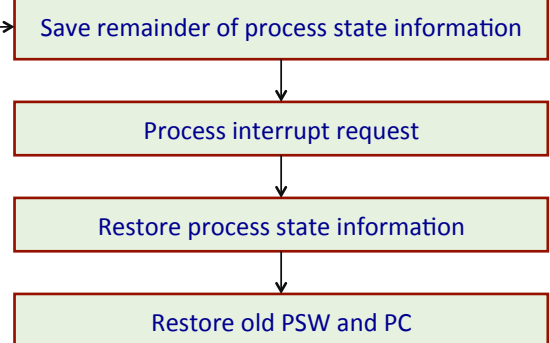
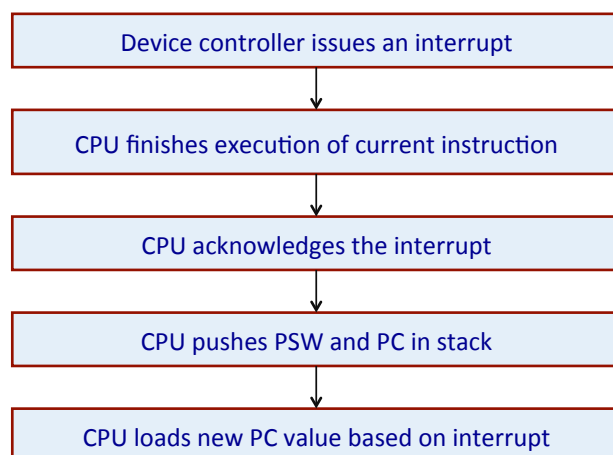
- At the end of the current instruction execution, the PC and *program status word (PSW)* are saved in stack automatically.
  - PSW contains status flags and other processor status information.
- The interrupt is acknowledged, the *interrupt vector* obtained, based on which control transfers to the appropriate ISS.
  - Different interrupting devices may have different ISS's.
- After handling the interrupt, the ISR executes a special *Return From Interrupt (RTI)* instruction.
  - Restores the PSW and returns control to the saved PC address.
  - Unlike normal RETURN where PSW is not restored.

- An instruction cycle typically consists of several machine cycles.

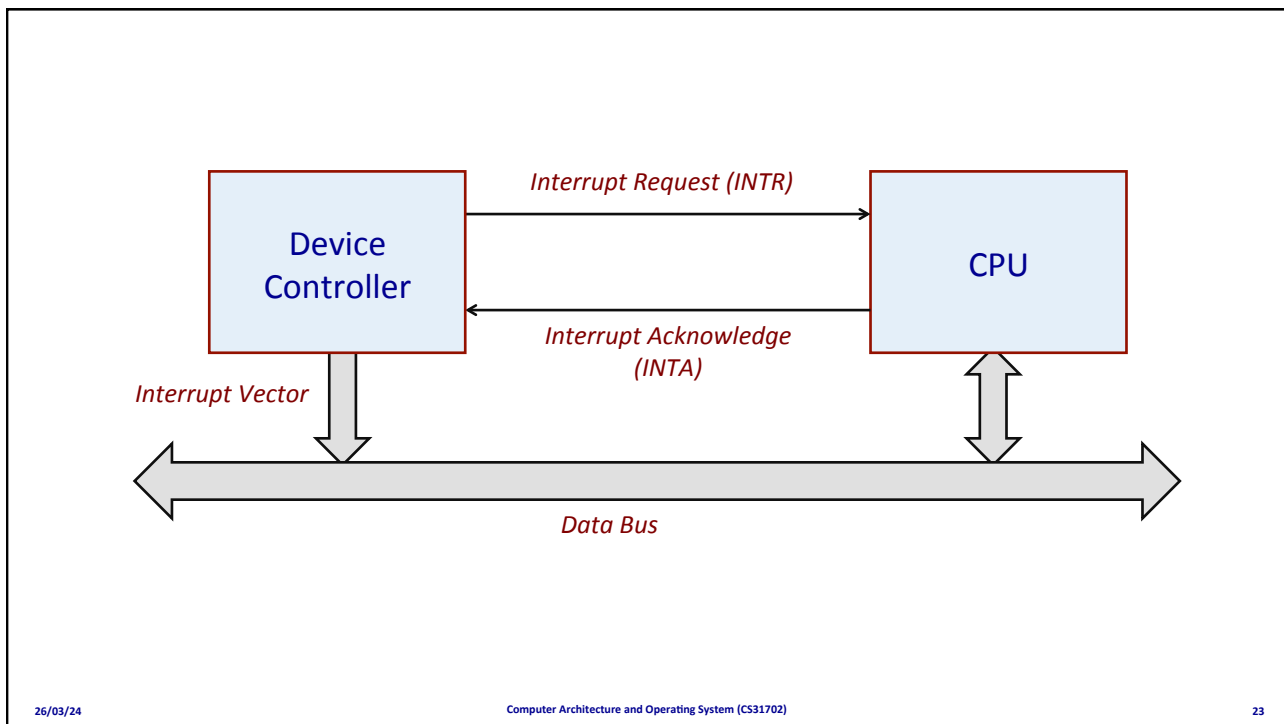


## General Interrupt Processing

### By hardware

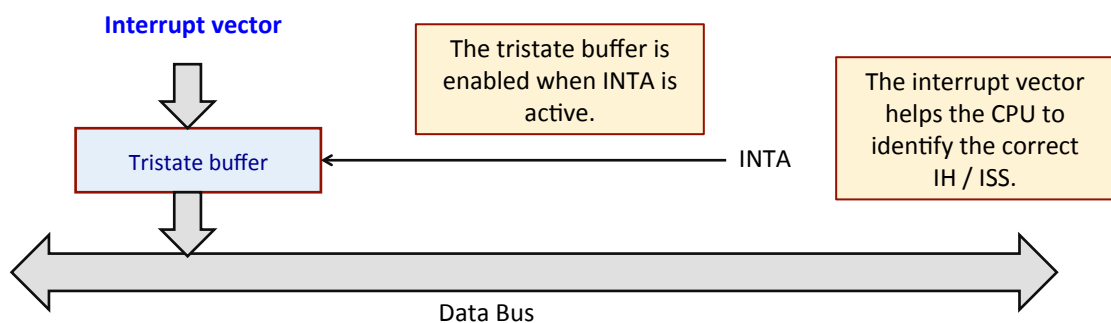


### By software




• How is the interrupt vector sent on the data bus in response to INTA?

- Device controller sends **INTR** to the CPU.
- CPU finishes the current instruction and sends back **INTA**.
- Device controller sends **interrupt vector** (or number) over data bus.
- CPU reads the interrupt vector, and identifies the device.



## Multiple Devices Interrupting the CPU

- A common solution is to use a *priority interrupt controller*.
  - The interrupt controller interacts with CPU on one side and multiple devices on the other side.
  - For simultaneous interrupt requests, interrupt priority is defined.
  - The interrupt controller is responsible for sending the interrupt vector to CPU.



26/03/24

Computer Architecture and Operating System (CS31702)

25

- How it works?
  - The *INTR* line is made active when one or more of the device(s) activate their interrupt request line.
$$INTR = INTR0 + INTR1 + INTR2 + INTR3$$
  - When the CPU sends back *INTA*, the interrupt controller sends back the corresponding acknowledge to the interrupting device, and puts the interrupt vector on the data bus.
  - The interrupt controller is programmable, where the interrupt vectors for the various interrupts can be programmed (specified).
  - For more than one interrupt request simultaneously active, a priority mechanism is used (e.g. *INTR0* is highest priority, followed by *INTR1*, etc.).

26/03/24

Computer Architecture and Operating System (CS31702)

26

## Cases that make interrupt handling difficult

- For some interrupts, it is not possible to finish the execution of the current instruction.
  - A special *RETURN* instruction is required that would return and restart the interrupted instructions.
- Some examples:
  - a) Page fault interrupt:** A memory location is being accessed that is not presently available in main memory.
  - b) Arithmetic exception:** Some error has occurred during some arithmetic operation (e.g. division by zero).

## Direct Memory Access (DMA)

## Introduction

- In the data transfer methods discussed under programmed I/O, it is assumed that machine instructions are used to transfer the data between I/O device and memory.
  - Not very suitable when large blocks of data are required to be transferred at high speed (e.g. transfer of a disk block).
- An alternate approach is *Direct Memory Access* (DMA).
  - Allows transfer of a block of data directly between an I/O device and memory, without continuous CPU intervention.

- Why programmed I/O is not suitable for high-speed data transfer?
  - a) Several program instructions have to be executed for each data word transferred between the I/O device and memory.
    - Suppose 20 instructions are required for each word transfer.
    - The CPI of the machine running at 1 GHz clock is 1.
    - So, 20 nsec is required for each word transfer → maximum 50 M words/sec
    - Data transfer rates of fast disks are higher than this figure.

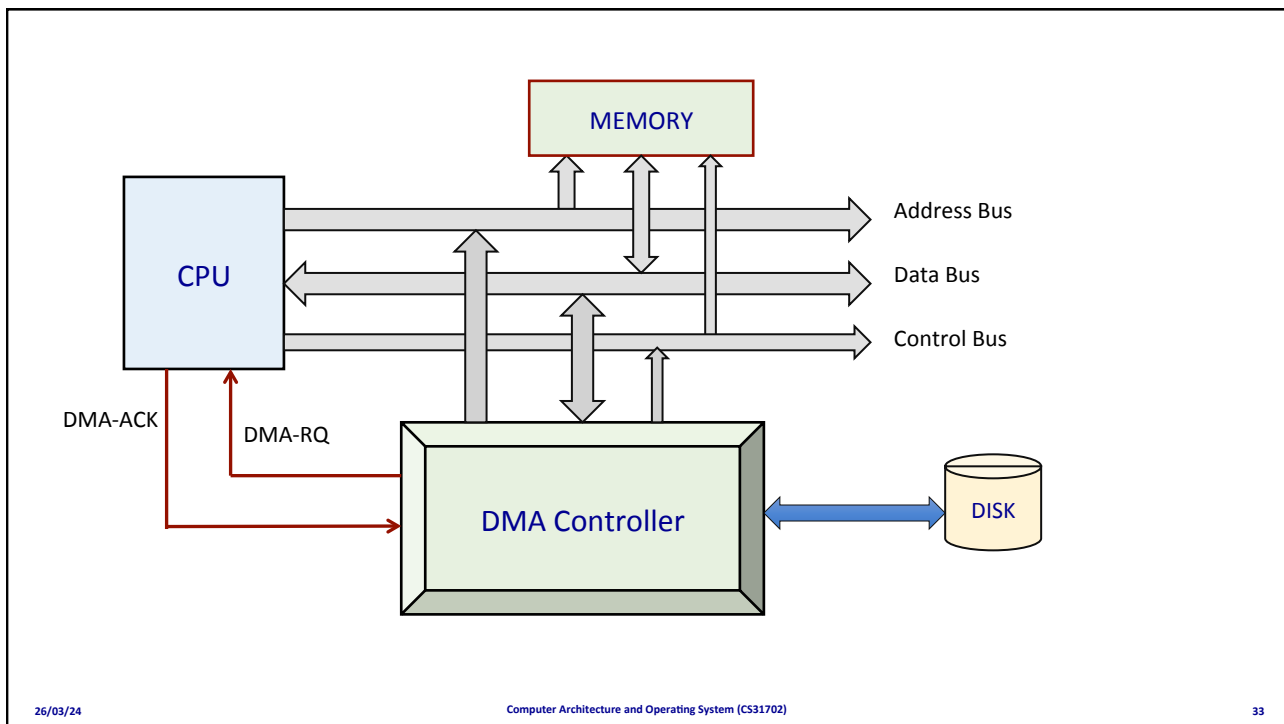
b) Many high speed peripheral devices like disk have a synchronous mode of operation, where data are transferred at a fixed rate.

- Consider a disk rotating at 7200 rpm, with rotational delay of 8.30 msec.
- Suppose there are 64 Kbytes of data recorded in every track.
- Once the disk head reaches the desired track, there will be a sustained data transfer at maximum rate  $64 \text{ Kbytes} / 8.30 \text{ msec} = 7.7 \text{ MB/s}$ .
- This sustained data transfer rate is comparable to the memory bandwidth, and cannot be handled by programmed I/O.

## DMA Controller

- A hardwired controller called the *DMA controller* can enable direct data transfer between I/O device (e.g. disk) and memory without CPU intervention.
  - No need to execute instructions to carry out data transfer.
  - Maximum data transfer speed will be determined by the rate with which memory read and write operations can be carried out.
  - Much faster than programmed I/O.






## Steps Involved

- When the CPU wants to transfer data, it initializes the DMA controller.
  - How many bytes to transfer, address in memory for the transfer.
- When the I/O device is ready for the transfer, the DMA controller sends **DMA-RQ** signal to the CPU.
- CPU waits till the next DMA breakpoint, relinquishes control of the bus (i.e. puts them in high impedance state), and sends **DMA-ACK** to DMA controller.
- Now DMA controller enables its bus interface, and transfers data directly to/from memory.
- When done, it deactivates the **DMA-RQ** signal.
- The CPU again begins to use the bus to access memory.

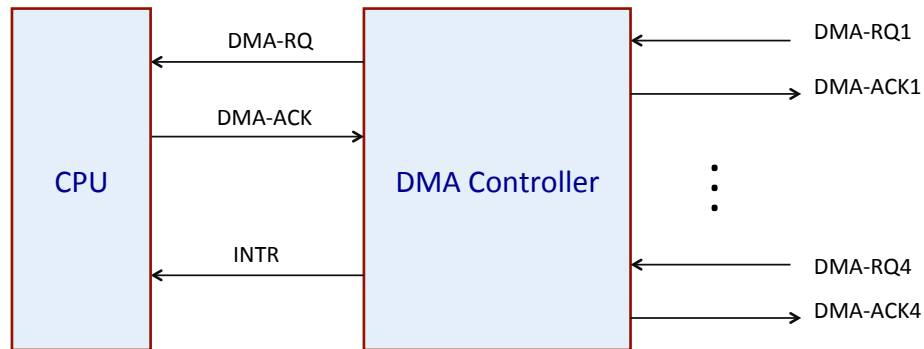
- The DMA breakpoints:
  - DMA request can be acknowledged at the end of any machine cycle.

*DMA breakpoints*



*Why cannot we have interrupt breakpoints at the end of any machine cycle?*

- For every DMA channel, the DMA controller will have three registers:
  - a) Memory address
  - b) Word count
  - c) Address of data on disk
- CPU initializes these registers before each DMA transfer operation.
- Before the data transfer, DMA controller requests the memory bus from the CPU.
- When the data transfer is complete, the DMA controller sends an interrupt signal to the CPU.



26/03/24

Computer Architecture and Operating System (CS31702)

37

## DMA Transfer Modes

- DMA transfer can take place in two modes:

### a) DMA cycle stealing

- The DMA controller requests for the for a few cycles 1 or 2.
- Preferably when the CPU is not using memory.
- DMA controller is said to steal cycles from the CPU without the CPU knowing it.

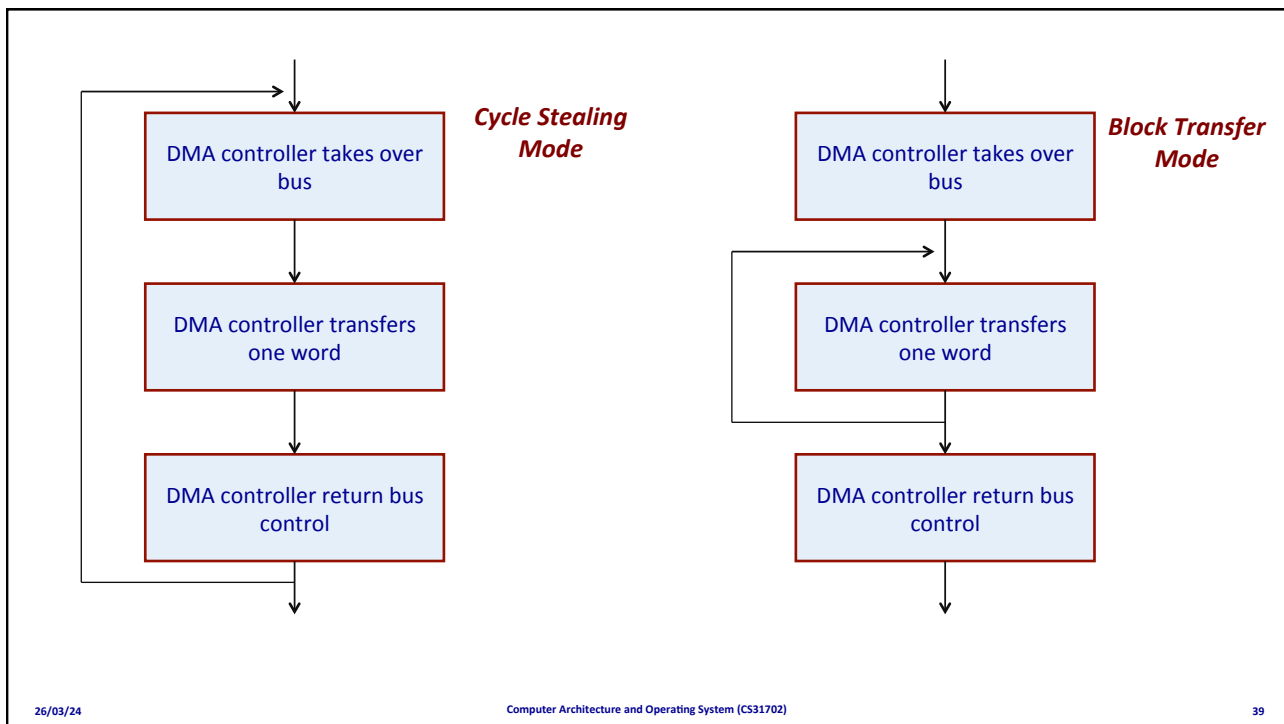
### b) DMA block transfer

- The DMA controller transfers the whole block of data without interruption.
- Results in maximum possible data transfer rate.
- CPU will lie idle during this period as it cannot fetch any instructions from memory.

26/03/24

Computer Architecture and Operating System (CS31702)

38



## Others Applications of DMA

- Other than data transfer to/from high-speed peripheral devices, DMA can be used in some other areas as well:
  - High-speed memory-to-memory block move.
  - Refreshing dynamic memory systems, by periodically generating dummy read requests to the columns.