

```
/*
```

PROCESS FORKING

In Linux we have multiple processes executing at the same time. This program illustrates the `fork()` system call which creates a new child process of the calling process. Go through the comments carefully before you compile and execute the program.

```
*/
```

```
#include <stdio.h>
```

```
#include <sys/ipc.h> /* You must include this file for all activities  
                      related to creation and interprocess communication.  
                      IPC stands for Inter-Process Communication */
```

```
/*
```

When the program begins executing from `main()`, it is a single process having a single process identifier (PID). When the `fork()` call is executed, a new process with a new PID is created. THE CODE OF THE NEW PROCESS IS IDENTICAL TO THE PARENT PROCESS, AND THE VALUES OF ALL VARIABLES IS A COPY OF THE VALUES OF THE VARIABLES OF THE PARENT AT THE TIME OF THE `fork()` CALL. The PIDs of the currently executing processes and the PIDs of their parents can be seen by using the "`ps -ef`" command from the command prompt.

```
*/
```

```
main()
```

```
{
```

```
    printf("Now there is a single process\n");  
    printf("Observe what happens when we fork\n");  
    printf("You can stop by pressing Control-c\n");  
    printf("Press Enter to continue... \n"); getchar();
```

```
    /*
```

For the parent process, the `fork()` system call returns the PID of the newly created child process. On the other hand the newly created child process, whose execution starts from the `fork()` statement (that is, its control flow starts from the `fork()` statement) receives the value 0 as return value from the `fork()` statement.

```
    */
```

```
    if (fork() == 0) {
```

```
        /*
```

We are here only if `fork()` returned zero. Therefore, this part of the code is executed only by the child process.

```
        */
```

```
        while (1)
```

```
            printf("\t\t\t Child executing\n ");
```

```
    }
```

```

        else {

            /*
             * We are here only if fork() returned a non-zero value.
             * Therefore this part of the code is executed only by
             * the original (or parent) process.
             */
            while (1)
                printf("Parent executing\n");
        }
    }
}

```

```

/*
After execution:

```

1. Did you observe how the parent and child take turns in execution?
 - Each process in the system gets a time slice to execute after which it is put away until its next turn. Time is maintained by a counter called "timer". At the end of a time slice, the timer interrupts the CPU, causing the CPU to execute an interrupt service routine. This routine calls the process scheduler which selects the next process to be executed in the next time slice.
2. Store the value returned by the fork() call in an integer variable and print out its value both in the parent as well as in the child. Does your observation match with your expectation?

```

*/

```

```

/*
PROCESS FORKING REVISITED

```

This program illustrates the fork() system call in more detail.

```

*/

```

```

#include <stdio.h>
#include <sys/ipc.h>

```

```

main()
{
    int i ;
    int x = 10 ;
    int pid1, pid2 ;

    printf("Before forking, the value of x is %d\n", x);

    /*
     * After forking, we make the parent and its two children
     * increment x in different ways to illustrate that they
     * have different copies of x
     */
}

```

```

*/
if ((pid1 = fork()) == 0) {

    /* First child process */
    for (i=0 ; i < 5; i++) {
        printf("\t\t\t At first child: x= %d\n", x);
        x= x+10;
        sleep(1) ; /* Sleep for 1 second */
    }
}
else {

    /* Parent process */

    /* Create another child process */
    if ((pid2 = fork()) == 0) {

        /* Second child process */
        for (i=0 ; i < 5; i++) {
            printf("\t\t\t\t\t At second child: x= %d\n", x);
            x= x+20;
            sleep(1) ; /* Sleep for 1 second */
        }
    }
    else {

        /* Parent process */
        for (i=0 ; i < 5; i++) {
            printf("At parent: x= %d\n", x);
            x= x+5;
            sleep(1) ; /* Sleep for 1 second */
        }

        /*
        The wait() system call causes the parent
        to wait for a child process to complete
        its execution. The input parameter can
        specify the PID of the child process for
        which it has to wait.
        */

        wait(pid1);
        wait(pid2);
    }
}
}
/*

```

SHARING MEMORY BETWEEN PROCESSES

In this example, we show how two processes can share a common portion of the memory. Recall that when a process forks, the

new child process has an identical copy of the variables of the parent process. After fork the parent and child can update their own copies of the variables in their own way, since they don't actually share the variable. Here we show how they can share memory, so that when one updates it, the other can see the change.

```
*/

#include <stdio.h>
#include <sys/ipc.h>
#include <sys/shm.h> /* This file is necessary for using shared
                      memory constructs
                      */

main()
{
    int shmid;
    int *a, *b;
    int i;

    /*
     The operating system keeps track of the set of shared memory
     segments. In order to acquire shared memory, we must first
     request the shared memory from the OS using the shmget()
     system call. The second parameter specifies the number of
     bytes of memory requested. shmget() returns a shared memory
     identifier (SHMID) which is an integer. Refer to the online
     man pages for details on the other two parameters of shmget()
    */
    shmid = shmget(IPC_PRIVATE, 2*sizeof(int), 0777|IPC_CREAT);
    /* We request an array of two integers */

    /*
     After forking, the parent and child must "attach" the shared
     memory to its local data segment. This is done by the shmat()
     system call. shmat() takes the SHMID of the shared memory
     segment as input parameter and returns the address at which
     the segment has been attached. Thus shmat() returns a char
     pointer.
    */
    if (fork() == 0) {

        /* Child Process */

        /* shmat() returns a char pointer which is typecast here
           to int and the address is stored in the int pointer b. */
        b = (int *) shmat(shmid, 0, 0);

        for( i=0; i< 10; i++) {
            sleep(1);
```

```

        printf("\t\t\t Child reads: %d,%d\n",b[0],b[1]);
    }

}
else {

    /* Parent Process */

    /* shmat() returns a char pointer which is typecast here
    to int and the address is stored in the int pointer a.
    Thus the memory locations a[0] and a[1] of the parent
    are the same as the memory locations b[0] and b[1] of
    the parent, since the memory is shared.
    */
    a = (int *) shmat(shmid, 0, 0);

    a[0] = 0; a[1] = 1;
    for( i=0; i< 10; i++) {
        sleep(1);
        a[0] = a[0] + a[1];
        a[1] = a[0] + a[1];
        printf("Parent writes: %d,%d\n",a[0],a[1]);
    }
    wait();
}
}

/*

```

POINTS TO NOTE:

In this case we find that the child reads all the values written by the parent. Also the child does not print the same values again.

1. Modify the sleep in the child process to sleep(2). What happens now?
2. Restore the sleep in the child process to sleep(1) and modify the sleep in the parent process to sleep(2). What happens now?

Thus we see that when the writer is faster than the reader, then the reader may miss some of the values written into the shared memory. Similarly, when the reader is faster than the writer, then the reader may read the same values more than once. Perfect inter-process communication requires synchronization between the reader and the writer.

Further note that "sleep" is not a synchronization construct. We use "sleep" to model some amount of computation which may exist in the process in a real world application.

*/

/*

COMMUNICATING VIA A PIPE

In the previous example we saw that communicating through shared memory requires synchronization between the two processes in order to ensure that all values that are written by one into the shared memory is read (and read once) by the other. Pipe is a system construct which facilitates such communication. A pipe is also a shared buffer. When a process tries to read from an empty pipe, it waits until someone has written something into the pipe. When the pipe is full, any process attempting to write into the pipe is made to wait.

*/

```
#include <stdio.h>
```

```
#include <unistd.h>    /* Include this file to use pipes */
```

```
#define BUFSIZE 80    /* We will write lines of 80 chars into the pipe.
                        The pipe has a large capacity and can accomodate
                        many such lines.
                        */
```

```
main()
```

```
{
```

```
    int fd[2];
```

```
    int n=0;
```

```
    int i;
```

```
    char line[BUFSIZE];
```

```
    /* A pipe is treated as a file by the system. You must have used
       fopen() to open a file. fopen() returns a "file pointer" which
       is used in fprintf(), fscanf(), fclose() etc. However, when we
       wish to perform reads and writes in blocks from a file, we use
       the system call "open" to open a file. Internally files are
       always opened using the "open" call. For each process the system
       maintains a "file descriptor table" (FDT) containing an entry
       for each file opened by that process. When a new file is opened,
       a new entry is created in the FDT, and the entry number is
       returned as an integer called "file descriptor".
```

Unlike in a file, we may want to both read and write from a pipe at the same time. Hence when a pipe is created, two file descriptors are created -- one for reading the pipe and one for writing into the pipe. The pipe() system call requires an array of two integers as parameter. The system returns the file descriptors through this array.

```
    */
```

```
    pipe(fd); /* fd[0] is for reading,
               fd[1] is for writing
```

```
    */
```

/* To illustrate the working of the pipe, we will make the child process write the integer n into the pipe and make the parent to read from the pipe. We put sleep in the writer process (in this case the child) to show that the reader process waits for the writer to write into the pipe.

To write a block of bytes into a pipe (or more generally into a file) the write() system call is used. Similarly read() is used to read a block of bytes from a file. Refer to the online man pages for these calls.

```

*/
if (fork() == 0) {

    close(fd[0]); /* The child will not read and
                  hence we close fd[0]
                  */

    for (i=0; i < 10; i++) {

        sprintf(line,"%d",n); /* Since write() accepts only
                               arrays of characters, we
                               first write the integer n
                               into the char array "line"
                               */
        write(fd[1], line, BUFSIZE);
        printf("Child writes: %d\n",n);
        n++;
        sleep(2);
    }
}
else {

    close(fd[1]); /* The parent will not write and
                  hence we close fd[1]
                  */

    for (i=0; i < 10; i++) {

        printf("\t\t\t Parent trying to read pipe\n");
        read(fd[0], line, BUFSIZE);
        sscanf(line,"%d",&n); /* Read the integer from the
                               line of characters read
                               from the pipe
                               */
        printf("\t\t\t Parent reads: %d\n",n);
    }
}
}
/*

```

In this program we illustrate the use of Berkeley sockets for interprocess communication across the network. We show the communication between a server process and a client process.

Since many server processes may be running in a system, we identify the desired server process by a "port number". Standard server processes have a worldwide unique port number associated with it. For example, the port number of SMTP (the sendmail process) is 25. To see a list of server processes and their port numbers see the file `/etc/services`

In this program, we choose port number 6000 for our server process. Here we shall demonstrate TCP connections only. For details and for other types of connections see:

Unix Network Programming
-- W. Richard Stevens, Prentice Hall India.

To create a TCP server process, we first need to open a "socket" using the `socket()` system call. This is similar to opening a file, and returns a socket descriptor. The socket is then bound to the desired port number. After this the process waits to "accept" client connections.

```
*/

#include <stdio.h>
#include <sys/types.h>

/* The following three files must be included for network programming */
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>

/* THE SERVER PROCESS */

/* Compile this program with cc server.c -o server
   and then execute it as ./server &
*/
main()
{
    int                sockfd, newsockfd ; /* Socket descriptors */
    int                clien;
    struct sockaddr_in cli_addr, serv_addr;

    int i;
    char buf[100];      /* We will use this buffer for communication */

    /* The following system call opens a socket. The first parameter
       indicates the family of the protocol to be followed. For internet
       protocols we use AF_INET. For TCP sockets the second parameter
       is SOCK_STREAM. The third parameter is set to 0 for user
```



```

    applications.
*/
if ((sockfd = socket(AF_INET, SOCK_STREAM, 0)) < 0) {
    printf("Cannot create socket\n");
    exit(0);
}

/* The structure "sockaddr_in" is defined in <netinet/in.h> for the
internet family of protocols. This has three main fields. The
field "sin_family" specifies the family and is therefore AF_INET
for the internet family. The field "sin_addr" specifies the
internet address of the server. This field is set to INADDR_ANY
for machines having a single IP address. The field "sin_port"
specifies the port number of the server.
*/
serv_addr.sin_family      = AF_INET;
serv_addr.sin_addr.s_addr = INADDR_ANY;
serv_addr.sin_port        = 6000;

/* With the information provided in serv_addr, we associate the server
with its port using the bind() system call.
*/
if (bind(sockfd, (struct sockaddr *) &serv_addr,
           sizeof(serv_addr)) < 0) {
    printf("Unable to bind local address\n");
    exit(0);
}

listen(sockfd, 5); /* This specifies that up to 5 concurrent client
                    requests will be queued up while the system is
                    executing the "accept" system call below.
                    */

/* In this program we are illustrating a concurrent server -- one
which forks to accept multiple client connections concurrently.
As soon as the server accepts a connection from a client, it
forks a child which communicates with the client, while the
parent becomes free to accept a new connection. To facilitate
this, the accept() system call returns a new socket descriptor
which can be used by the child. The parent continues with the
original socket descriptor.
*/
while (1) {

    /* The accept() system call accepts a client connection.
    It blocks the server until a client request comes.

    The accept() system call fills up the client's details
    in a struct sockaddr which is passed as a parameter.
    The length of the structure is noted in clien. Note
    that the new socket descriptor returned by the accept()

```

```

        system call is stored in "newsockfd".
    */
    clilen = sizeof(cli_addr);
    newsockfd = accept(sockfd, (struct sockaddr *) &cli_addr,
                      &clilen) ;

    if (newsockfd < 0) {
        printf("Accept error\n");
        exit(0);
    }

    /* Having successfully accepted a client connection, the
       server now forks. The parent closes the new socket
       descriptor and loops back to accept the next connection.
    */
    if (fork() == 0) {

        /* This child process will now communicate with the
           client through the send() and recv() system calls.
        */
        close(sockfd); /* Close the old socket since all
                       communications will be through
                       the new socket.
        */

        /* We initialize the buffer, copy the message to it,
           and send the message to the client.
        */
        for(i=0; i < 100; i++) buf[i] = '\0';
        strcpy(buf, "Message from server");
        send(newsockfd, buf, 100, 0);

        /* We again initialize the buffer, and receive a
           message from the client.
        */
        for(i=0; i < 100; i++) buf[i] = '\0';
        recv(newsockfd, buf, 100, 0);
        printf("%s\n", buf);

        close(newsockfd);
        exit(0);
    }

    close(newsockfd);
}
}

```

```

/*

```

THE CLIENT PROCESS

Please read the file server.c before you read this file. To run this,

you must first change the IP address specified in the line:

```
serv_addr.sin_addr.s_addr = inet_addr("144.16.202.221");
```

to the IP-address of the machine where you are running the server.

```
*/
```

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
```

```
main()
{
```

```
    int                sockfd ;
    struct sockaddr_in  serv_addr;
```

```
    int i;
    char buf[100];
```

```
    /* Opening a socket is exactly similar to the server process */
    if ((sockfd = socket(AF_INET, SOCK_STREAM, 0)) < 0) {
        printf("Unable to create socket\n");
        exit(0);
    }
```

```
    /* Recall that we specified INADDR_ANY when we specified the server
       address in the server. Since the client can run on a different
       machine, we must specify the IP address of the server.
```

TO RUN THIS CLIENT, YOU MUST CHANGE THE IP ADDRESS SPECIFIED
BELOW TO THE IP ADDRESS OF THE MACHINE WHERE YOU ARE RUNNING
THE SERVER.

```
    */
    serv_addr.sin_family      = AF_INET;
    serv_addr.sin_addr.s_addr = inet_addr("144.16.202.221");
    serv_addr.sin_port        = 6000;
```

```
    /* With the information specified in serv_addr, the connect()
       system call establishes a connection with the server process.
```

```
    */
    if ((connect(sockfd, (struct sockaddr *) &serv_addr,
                  sizeof(serv_addr))) < 0) {
        printf("Unable to connect to server\n");
        exit(0);
    }
```

```
    /* After connection, the client can send or receive messages.
       However, please note that recv() will block when the
       server is not sending and vice versa. Similarly send() will
```

block when the server is not receiving and vice versa. For non-blocking modes, refer to the online man pages.

```
*/  
for(i=0; i < 100; i++) buf[i] = '\0';  
recv(sockfd, buf, 100, 0);  
printf("%s\n", buf);  
  
for(i=0; i < 100; i++) buf[i] = '\0';  
strcpy(buf, "Message from client");  
send(sockfd, buf, 100, 0);  
  
close(sockfd);  
}
```
